

Project 1A: SUMMA Algorithm

David Noble, Jack Chua

September 24, 2011

1 Overview

The objective of this checkpoint is to implement the SUMMA distributed matrix multiplication algorithm for non-square matrices and analyze the results.

2 Plots

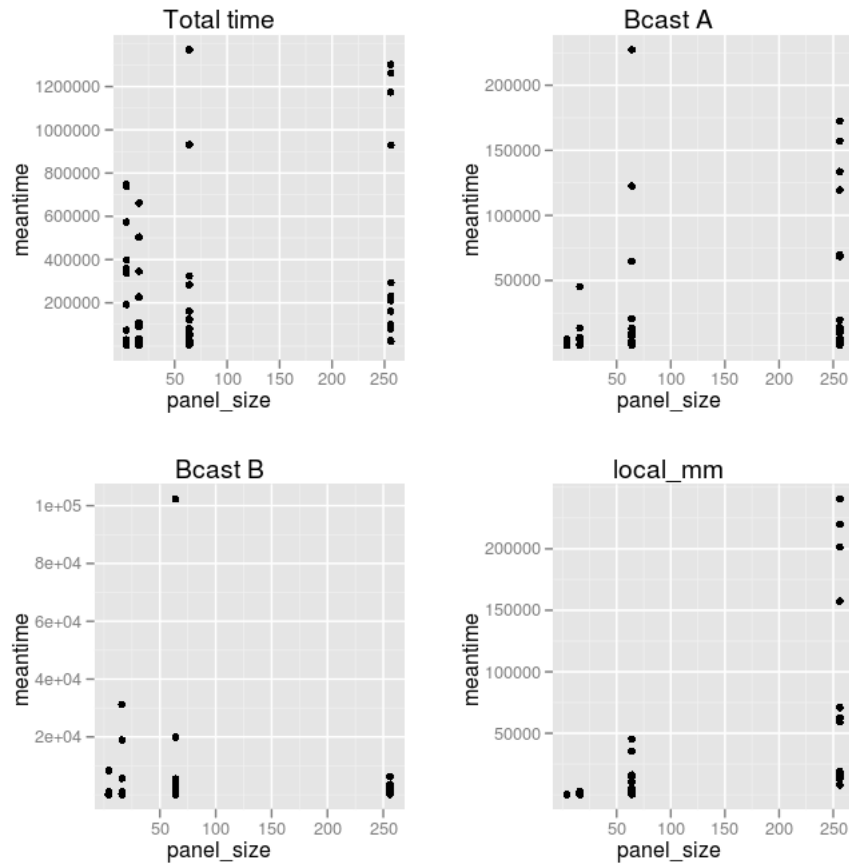


Figure 1: How total time, broadcast time, and local multiplication time varies as panel size varies

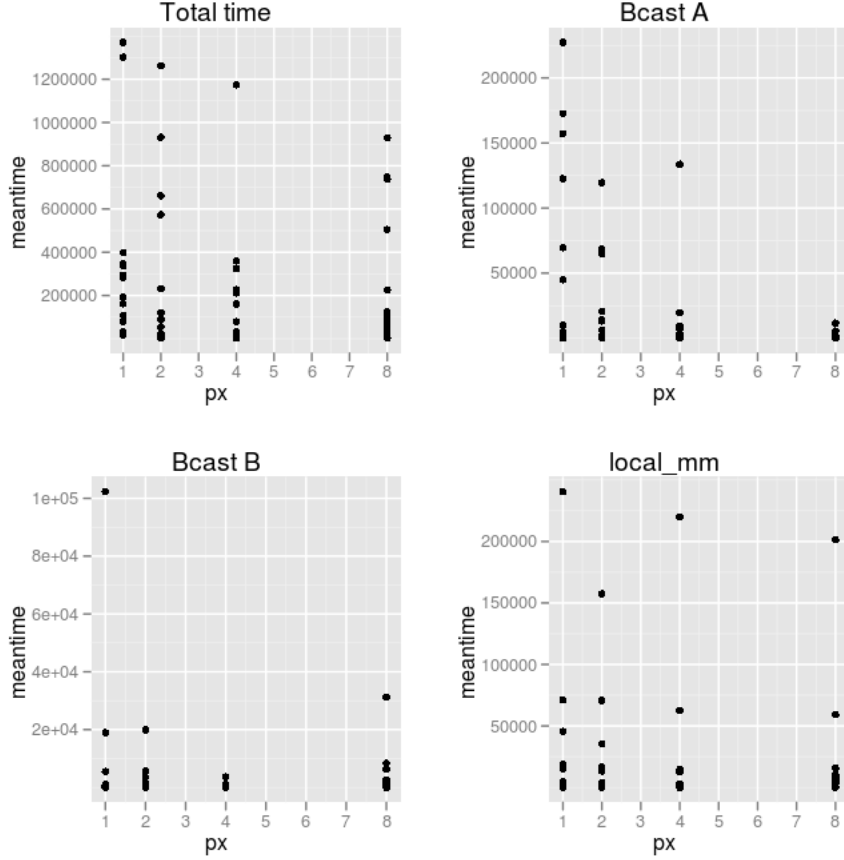


Figure 2: How total time, broadcast time, and local multiplication time varies according to grid layout

3 Analysis

Methodology. These plots were generated by taking the average times collected from each slice for every unique test case. Uniqueness is defined by the combination of variables m , n , k , px , py , and $panel_size$. We collected the total time elapsed, the time taken to broadcast Ablock and Bblock, and the time taken to compute the local matrix multiply. We then used R to plot the times against panel size and grid layout.

Analysis of panel size. We analyzed how changing the panel size might affect various potential bottlenecks in the algorithm.

- Looking at the **total time elapsed**, we noticed that increasing the panel size tended to **worsen** performance. This makes sense because in a broad sense, we shouldn't achieve the benefits of parallelism if the panel size is too big and we get closer and closer to a (by-nature) sequential outer product.
- The **time taken to broadcast Ablock** seemed to **benefit from smaller panel sizes**. The largest panel size provided the largest concentration of high time values, while the smallest panel size provided the largest concentration of small time values.
- The **time taken to broadcast Bblock** seemed to **benefit from the largest panel size, and the smallest panel size**. The middling panel sizes seemed to provide the worst performance.
- Finally, **time taken for local matrix multiplication** showed a clear increase as the panel size grew. This makes obvious sense as larger blocks call for larger overhead during sequential multiplication.

A feasible conclusion is that the smaller the panel size is, the better - this allows use to benefit from parallelism as much as possible. However, it does not seem like the **absolute smallest** panel size is the optimal configuration. There may be some small $panel_size$ between 1 and 5 that is optimal from an $\alpha - \beta$ perspective.

Analysis of process grid layout. We then analyzed how changing the grid layout might affect the times.

- Over all of our test cases (which include all panel sizes and matrix dimensions), using a square layout improved the total running time of the algorithm. We have enough observations to say fairly confidently that this improvement is not luck - a square layout probably lends to the best tradeoff between latency and bandwidth.
- Along the same lines, we can see that our broadcast times are improved with a square layout. This probably contributed the most to total running time.
- The square layout again seemed to help the local matrix multiply, but not as evidently.

Best configurations. We found the best configurations via the following calls:

```
> ttimedata[ttimedata$meantime == min(ttimedata$meantime),]
m    n    k    px    py panel_size meantime
256 256 256    4   16     16      2354
> Atimedata[Atimedata$meantime == min(Atimedata$meantime),]

m    n    k    px    py panel_size meantime
256 256 1024  8     8       6.453125
> Btimedata[Btimedata$meantime == min(Btimedata$meantime),]
m    n    k    px    py panel_size meantime
256 256 1024  1     64      0.1171875
> mtimedata[mtimedata$meantime == min(mtimedata$meantime),]

m    n    k    px    py panel_size meantime
256 256 256    4     16      7.40625
```

A few observations:

- As expected, passing in the smallest matrices possible resulted in the smallest (mean) total time elapsed.
- For some odd reason, the 1×64 grid layout provided the best (mean) broadcast time for Bblock.
- The 8×8 grid layout provided the best (mean) broadcast time for Ablock.

4 Performance Model

We generalize some results from the square SUMMA model, which states

$$T_{net} = \alpha \frac{n}{b} \log(p) + \frac{n^2}{\beta} \frac{\log(p)}{\sqrt{p}}$$

We didn't have time to really derive a specific formula for rectangular matrices, but perhaps a feasible way to look at it is to view $n = \text{average}(m, n, k)$. Since n is a marker for the size of a matrix, doing this might not be too far off from the truth.

5 Conclusions

Our analysis shows that a small panel size relative to block size and a square grid layout can definitively improve matrix multiplication over a naive implementation.